

Mindy — AI Capability for the DLA Enterprise / Agentic AI Requirement

Prepared for: Servexo Protective Services — GSA MRAS RFI, "OSW — AI Vendor Market Research" (DLA Enterprise Artificial Intelligence & Agentic AI Rapid Support Services) **Teaming partner:** GovCon Giants / Mindy (AI technical capability) · getmindy.ai · eric@govcongiants.com · 786-208-2071

Who we are

Mindy is a production AI platform delivering **grounded large-language-model (LLM) and retrieval-augmented generation (RAG)** over authoritative federal data — 317,000+ contractor records and 88,000+ live federal solicitations sourced from USASpending and SAM.gov — with **citation-backed answers, hallucination controls, and bounded agentic workflows** for market research, incumbent analysis, and response drafting.

Operating history: federal market-intelligence software for the government-contracting community **since 2024** (product lineage: OpenGov IQ → GovCon Giants marketplace → Market Assassin → Mindy). **Current scale:** ~1,900 active users/customers — ~1,400 contractors receiving daily AI opportunity alerts and ~480 on **Mindy Pro** (AI-generated market briefings). Commercial product per FAR Part 12. **Capability website:** getmindy.ai

On this requirement, **Servexo provides the GSA Schedule vehicle, cleared staffing, and program delivery; Mindy provides the AI/LLM/RAG and agentic-automation technical capability and federal domain expertise** the requirement centers on.

1. Production LLM/RAG integrated with authoritative data + legacy systems

Our architecture delivers a secure, production-grade LLM/RAG capability — not a prototype:

- **Ingestion & integration:** connectors to authoritative and legacy data sources; normalized, NUL-safe ingestion; semantic chunking with structured metadata (agency, NAICS/PSC, date, document type).
- **Vector store & retrieval:** embeddings to a managed vector store; **metadata-filtered retrieval** so answers draw only from the authorized, relevant corpus.
- **Grounding & citations:** every answer carries **inline citations to the source record and as-of date**. Responses that cannot be grounded **abstain** rather than fabricate.
- **Hallucination reduction:** retrieval-only context, citation enforcement, confidence gating, abstain-on-insufficient-evidence, and an LLM-as-judge fact-check on high-stakes output.
- **Model evaluation:** labeled gold dataset; precision/recall on retrieval and factuality scoring, with regression testing on every change.
- **Security & sustainment:** least-privilege service accounts, encryption in transit and at rest, full audit logging, a provider-agnostic model fallback chain (no single-vendor lock-in), and cost controls — engineered to run continuously in production.

2. AI-enabled workflow automation with agentic controls

We match the automation method to the task and enforce strict governance on autonomous action:

- **Method selection:** rules-based automation for deterministic steps; AI-assisted (human-in-the-loop) for drafting; retrieval-enabled for grounded answers; **agentic execution only where multi-step autonomy genuinely earns its risk**.
- **Human oversight & control:** approval routing, **bounded/scoped permissions**, reversible operations with rollback, full **audit logging and traceability** on every agentic action.
- **Safeguards:** defenses against unauthorized actions, invalid tool calls, prompt injection, and data leakage — via tool allow-listing, input/output filtering, and least-privilege so an agent cannot act outside its lane.

3. Delivery & staffing

Servexo leads cleared staffing and program delivery (Secret / IT-II personnel across AI/ML, LLM engineering, automation, integration, testing, responsible AI, cybersecurity, and project management), with a phased prototype → production → sustainment model over the first 12 months. Mindy defines the AI/ML and LLM-engineering role requirements and provides technical onboarding so cleared staff execute to standard.

Relevant experience (commercial — FAR Part 12)

Mindy (getmindy.ai) — a deployed, commercial AI platform delivering grounded LLM/RAG and bounded agentic workflows over authoritative federal data. In production since 2024 (OpenGov-IQ-onward lineage); ~1,900 active users/customers today. Demonstrates exactly the capabilities this requirement calls for: secure enterprise LLM/RAG integrated with authoritative data,

citation-backed grounding, hallucination controls, and governed agentic automation. *Commercial item — no CPARS.* Capability URL: getmindy.ai.

GovCon Giants / Mindy — AI technical teaming partner. Contact: Eric Coffie · eric@govcongiants.com · 786-208-2071 · getmindy.ai